

# CS4.301 Data & Applications

Ponnurangam Kumaraguru ("PK")  
#ProfGiri @ IIIT Hyderabad



pk.profgiri



/in/ponguru



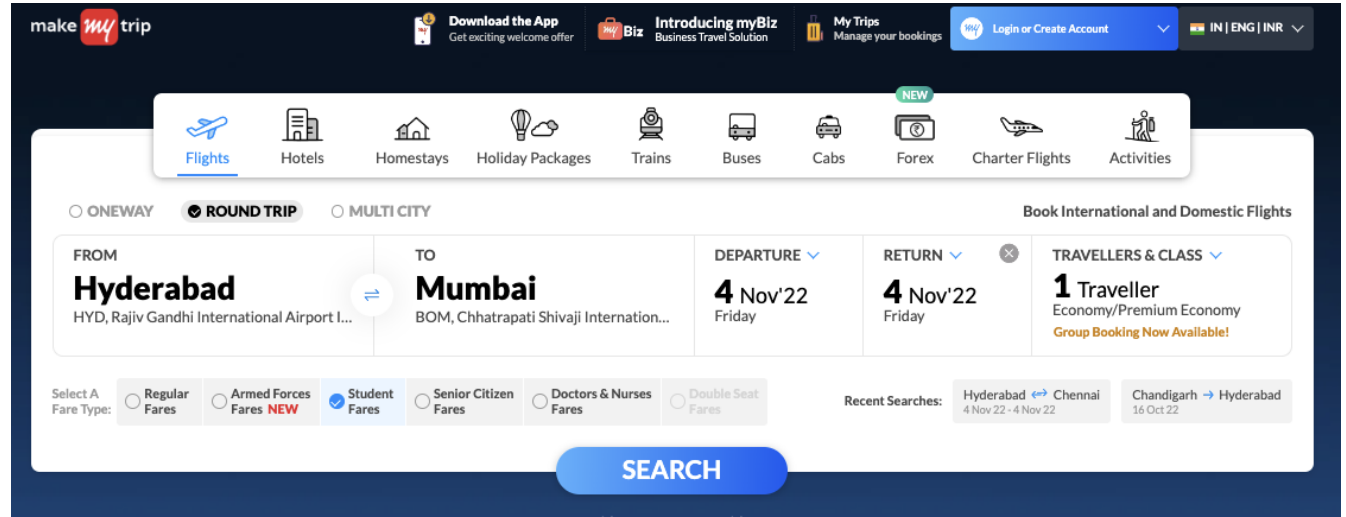
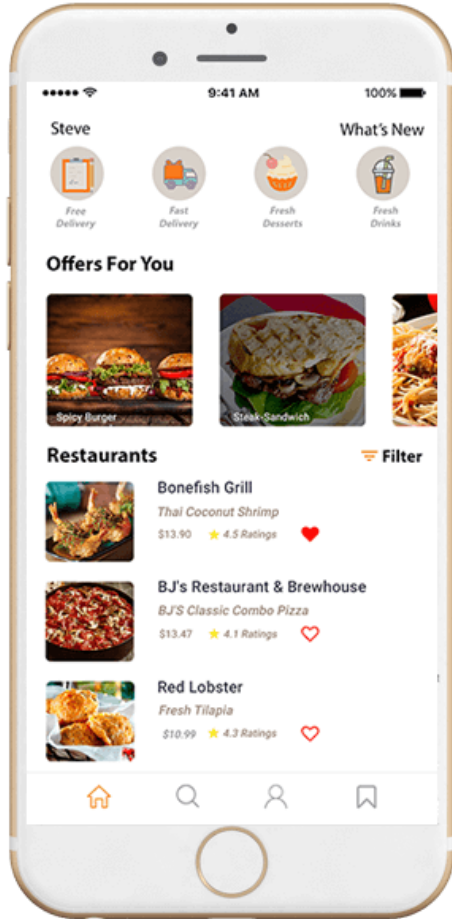
@ponguru




Ponnurangam.kumaraguru

# Protocol





What kind of data are we generating here?



What kind of data are we generating here?

# Grading, Relative

Type of Evaluation	Weightage (in %)
Class Quizzes (3)	15
Assignments / Homeworks (4)	20
Mid sem exam (Quiz-2 as scheduled in almanac)	15
Project	30
End Sem Exam	20

# TAs

12 TAs 13

Students will be assigned among TAs for all evaluations

Kaathik  
Goyal

Aditya Mishra

Anika Roy

Anirudh Vempati

Anish R Joishy

Chetan Mahipal

Debangana Mishra

Devika Umesh Bej

Jain Hemang Ashok

Pratishtha Saxena

Priyanshul Govil

Rohan Kumar

Shailender Goyal

Tejas Cavale

Memory challenge for me 😊

# Basic Definitions

## Database:

A collection of related data.

## Data:

Known facts that can be recorded and have an implicit meaning.

## Mini-world:

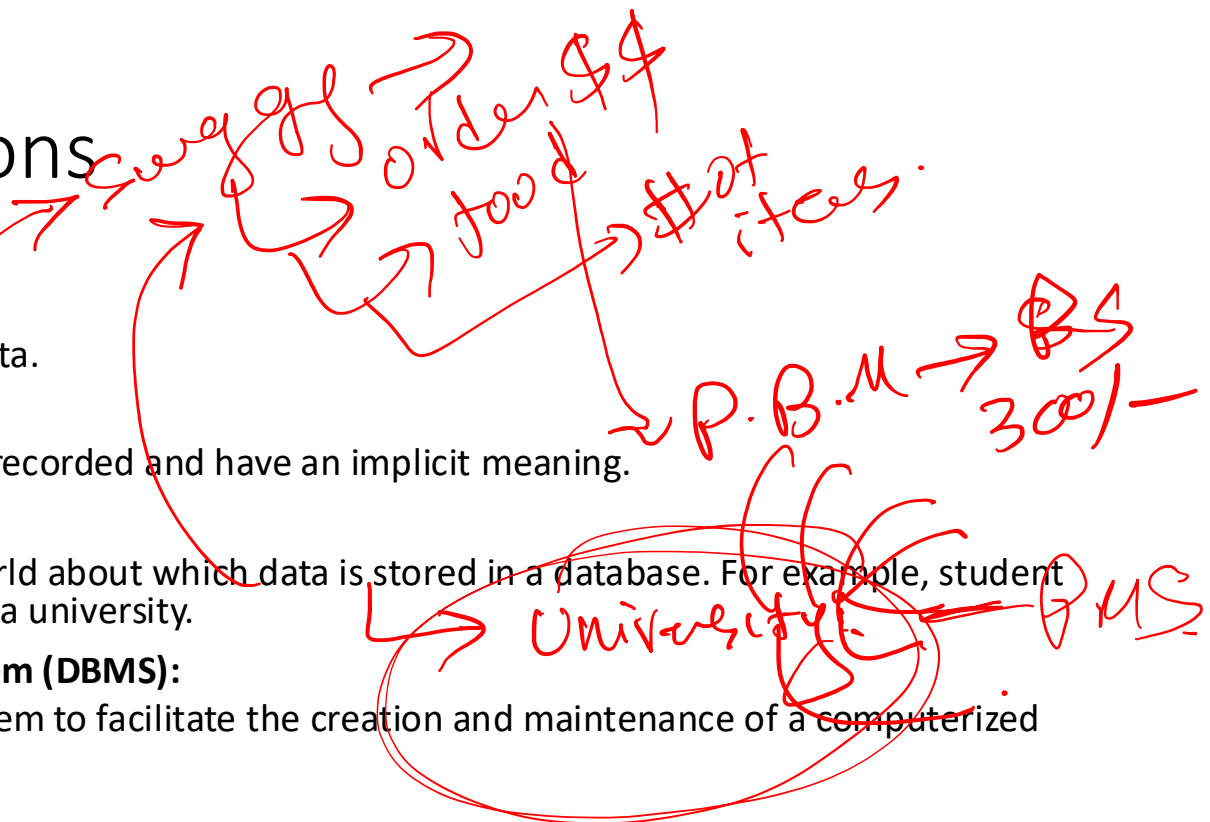
Some part of the real world about which data is stored in a database. For example, student grades and transcripts at a university.

## Database Management System (DBMS):

A software package/ system to facilitate the creation and maintenance of a computerized database.

## Database System:

The DBMS software together with the data itself. Sometimes, the applications are also included.



# What is a Database?

Data: factual (undoubted) information that can be recorded and have implicit meaning

University → Class / Lab → DNA - 243 - 137AS -  
↓ Plan -  
Notes.

A database is a collection of related data

— User —  
— Parents —  
— —



# What is a Database?

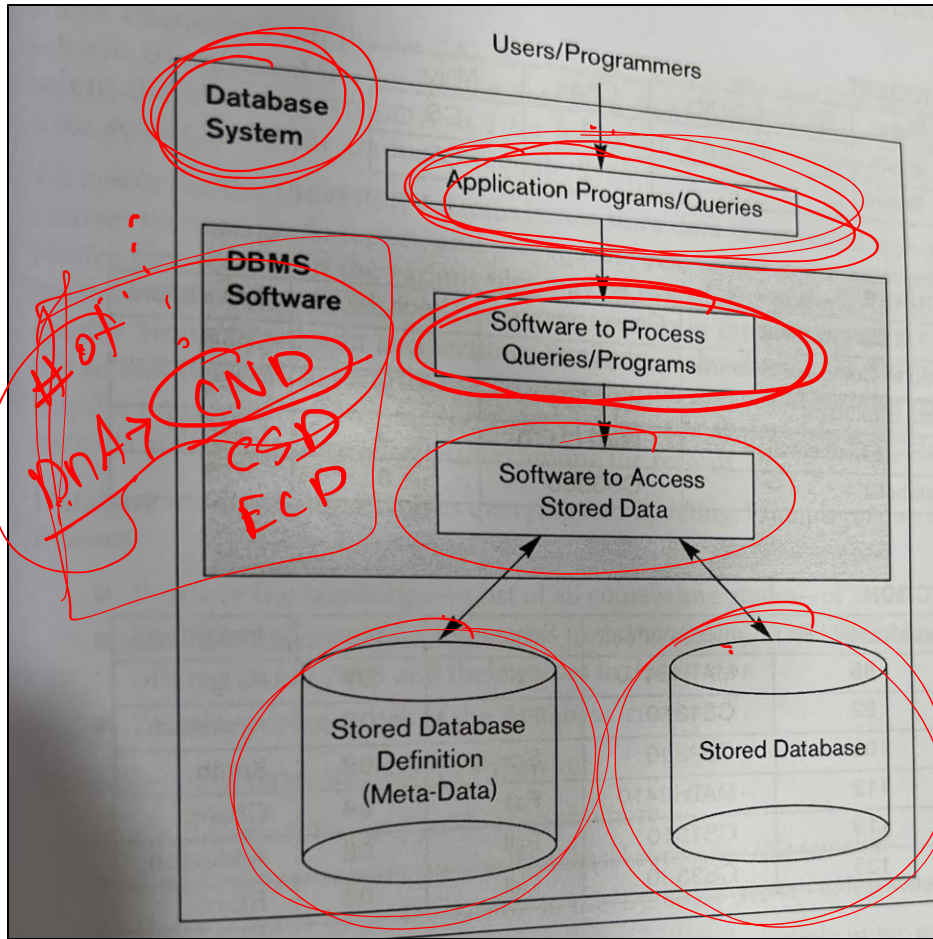
A database has the following implicit properties:

A database represents some aspect of the real world (mini-world or Universe of Discourse (UoD))

A database is a logically coherent (associated, related) collection of data with some inherent meaning

A database is designed, built and populated with data for a specific purpose

It has an intended group of users and some preconceived (already thought of) applications in which these users are interested



Book

# Simplified Database

TA hrs  
my hrs

A. - CND  
u3 - CND

NL -> MU Q...

# Example of Database: University

Data:

STUDENTs

COURSEs

SECTIONs (of COURSEs)

(academic) DEPARTMENTs

INSTRUCTORs

Relation:

SECTIONs are of specific COURSEs

STUDENTs take SECTIONs

COURSEs have prerequisite COURSEs

INSTRUCTORs teach SECTIONs

COURSEs are offered by DEPARTMENTs

STUDENTs major in DEPARTMENTs

AF → FE  
PX → US  
CNP

# Database

definitions  
IMS + master

STUDENT

Name	Student Number	Class	Major
Smith	17	1	COSC
Brown	8	2	COSC

GRADE REPORT

Student Number	Section- Identifier	Grade
17	85	A
18	102	B+

PREREQUISITE

Course Number	Prerequisite Number
COSC3380	COSC3320
COSC3320	COSC1310

COURSE

Course Name	Course Number	Credit Hours	Department
Intro to CS	COSC1310	4	COSC
Data Structures	COSC3320	4	COSC
Discrete Mathematics	MATH2410	3	MATH
Data Base	COSC3380	3	COSC

SECTION

Section- Identifier	Course Number	Semester	Year	Instructor
85	MATH2410	Fall	91	King
92	COSC1310	Fall	91	Anderson
102	COSC3320	Spring	92	Knuth
135	COSC3380	Fall	92	Stone

DL SNA2

# Database catalogue

## RELATIONS

Relation_name	No_of_columns
STUDENT	4
COURSE	4
SECTION	5
GRADE_REPORT	3
PREREQUISITE	2

## COLUMNS

Column_name	Data_type	Belongs_to_relation
Name	Character (30)	STUDENT
Student_number	Character (4)	STUDENT
Class	Integer (1)	STUDENT
Major	Major_type	STUDENT
Course_name	Character (10)	COURSE
Course_number	XXXXNNNN	COURSE
....	....	....
....	....	....
....	....	....
Prerequisite_number	XXXXNNNN	PREREQUISITE

Note: Major\_type is defined as an enumerated type with all known majors.  
XXXXNNNN is used to define a type with four alphabetic characters followed by four

# Views

Many users to DB

Each users may require a different view

View may be a subset or virtual data derived

IMS DB.

$D(A)$  Dy. A. Gled  $D(R)$  D. V. K. Shore.

(a)

		Student_transcript			
Student_name	Course_number	Grade	Semester	Year	Section_id
Smith	CS1310	C	Fall	08	119
	MATH2410	B	Fall	08	112
	MATH2410	A	Fall	07	85
Brown	MATH2410	A	Fall	07	92
	CS1310	A	Fall	08	102
	CS3320	B	Spring	08	135
	CS3380	A	Fall	08	

(b)

Course_name	Course_number	Prerequisites
Database	CS3380	CS3320
		MATH2410
Data Structures	CS3320	CS1310

**Figure 1.5**

Two views derived from the database in Figure 1.2. (a) The TRANSCRIPT view. (b) The COURSE\_PREREQUISITES view.

This lecture

# No replies this year 😞



**Re: Good morning.. Good luck with the course....**

by [Mayaank Ashok](#) - Monday, 25 September 2023, 4:51 AM

Looking forward to see you too sir :)



**Re: Good morning.. Good luck with the course....**

by [Kavish Kapoor](#) - Monday, 25 September 2023, 4:52 AM

Looking forward to see you too sir.



**Re: Good morning.. Good luck with the course....**

by [Maneesh Manoj](#) - Monday, 25 September 2023, 4:54 AM

Eagerly waiting your lectures sir :)



**Re: Good morning.. Good luck with the course....**

by [Nitin Avuthu](#) - Monday, 25 September 2023, 4:58 AM

Eager for tomorrow sir



**Re: Good morning.. Good luck with the course....**

by [Madhav Jayachandran Ramachandran](#) - Monday, 25 September 2023, 5:24 AM

I am eagerly looking forward to attending your class, sir.






# Online Transaction Processing (OLTP)

Multiuser DB

Concurrency control

Flight ticket booking, seats available

# Transaction



Executing program or process that includes one or more database accesses, reading or updating of database records

## Properties [ACID]

Atomicity: either all are executed or none are executed [A/c A (Read, Write) → A/c B (Read, Write)]

Consistency: any data written to a DB must be valid according to the defined rules [telephone number]

Isolation: each transaction appears to execute in isolation, even though 100s may be executing at the same time [updating the seat preference]

Durability: guarantees that once a transaction has been committed, it will remain committed even in the case of a system failure

# Actors on the Scene: Day-to-Day use of DB

## Database administrators

authorizing access to DB, coordinating & monitoring its use, accountable for security breaches & response time

## Database designers

responsible for identifying the data to be stored in the DB, interact with potential group of users and develop *views* of the DB

## End Users: Casual, naïve / parametric, sophisticated, stand-alone users

Casual: occasional users, typically middle or high-level managers

Naïve / parametric: constantly updating the db using *canned transaction*, done using mobile apps

bank tellers checking balances post withdrawals & deposits

reservation agents checking for availability

social media users post and read items on platforms

# Actors on the Scene: Day-to-Day use of DB

End Users: Casual, naïve, sophisticated, stand-alone users

sophisticated: thoroughly familiarize themselves with all facilities of DBMS, implement their own, complex requirements

stand-alone: maintain personal DB using ready-made programs;  
TALLY

System analysts & application programmers

determine the requirements of end-users, including naïve, develop specifications for canned transactions

map implement above specifications as programs, they test – debug  
maintain these canned transactions

software developers / engineers play these roles sometimes

# Actors Behind the Scene: Maintain the DB

## DBMS designers & implementers

design and implement the DBMS modules; complex modules like query language processing, interface processing, controlling concurrency, handling data recovery & security

## Tool developers

design & implement tools; optional packages that are often purchased separately; facilitate DB modeling & design, system design, and improved performance

## Operators & maintenance personnel

responsible for running & maintenance of the hardware & software environment for DB

# Advantages of using DBMS approach

## Controlling redundancy

- redundancy in storing the same data multiple times

- e.g. student details in university maintained by acad & finance office separately

- duplication of efforts, storage space, inconsistent data [Jan-19-1998 vs Jan-29-1998]

- ideally student details in only one place, *data normalization*

- keeping all needed data together, *denormalization*

# Advantages of using DBMS approach

Restricting unauthorized access

your grades accessible to only some; my salary and personal details only to some [hopefully 😊]

Providing storage structures and search techniques for efficient query processing

efficiently executing queries & updates; creating *indexes* and maintaining it; *buffering* & *caching* modules



# Advantages of using DBMS approach

## Providing backup & recovery

- provide facilities for recovering from hardware & software failures

- complex updates, should not crash; if crash what state to recover

## Providing multiple user interfaces

- apps for mobile users; query language for causal; programming language for application programmers; forms / command codes for parametric; menu & natural language interfaces for standalone

# Advantages of using DBMS approach

## Representing Complex relationship among data

DBMS must have the capability to represent a variety of complex relationships among the data, to define new relationships as they arise, and to retrieve / update related data easily & efficiently

## Enforcing integrity constraints

student name: 30 alphabetic characters; record in one file must be related to records in other files [e.g. every SECTION record must be related to a COURSE record] *referential integrity*

uniqueness on data item values [e.g. every COURSE record must have a unique value for COURSE\_NUMBER] *key or uniqueness constraint*

# Advantages of using DBMS approach

Permitting inferencing and actions using rules and triggers

triggers associated with tables; trigger is a rule activated by updates to the table results in performing some addition operations to other tables, sending messages, etc.

stored procedures are invoked appropriately when some conditions are met

# Historical Development of Database Technology

## Early Database Applications:

The Hierarchical and Network Models were introduced in mid 1960s and dominated during the seventies.

A bulk of the worldwide database processing still occurs using these models, particularly, the hierarchical model using IBM's IMS system.

## Relational Model based Systems:

Relational model was originally introduced in 1970, was heavily researched and experimented within IBM Research and several universities.

Relational DBMS Products emerged in the early 1980s.

# Historical Development of Database Technology (continued)

## Object-oriented and emerging applications:

Object-Oriented Database Management Systems (OODBMSs) were introduced in late 1980s and early 1990s to cater to the need of complex data processing in CAD and other applications.

Their use has not taken off much.

Many relational DBMSs have incorporated object database concepts, leading to a new category called *object-relational* DBMSs (ORDBMSs)

*Extended relational* systems add further capabilities (e.g. for multimedia data, text, XML, and other data types)

# Historical Development of Database Technology (continued)

## Data on the Web and E-commerce Applications:

Web contains data in HTML (Hypertext markup language) with links among pages.

This has given rise to a new set of applications and E-commerce is using new standards like XML (eXtended Markup Language). (see Ch. 13).

Script programming languages such as PHP and JavaScript allow generation of dynamic Web pages that are partially generated from a database (see Ch. 11).

Also allow database updates through Web pages

# Extending Database Capabilities (1)

New functionality is being added to DBMSs in the following areas:

Scientific Applications – Physics, Chemistry, Biology - Genetics

Earth and Atmospheric Sciences and Astronomy

XML (eXtensible Markup Language)

Image Storage and Management

Audio and Video Data Management

Data Warehousing and Data Mining – a very major area for future development using new technologies (see Chapters 28-29)

~~Spatial Data Management and Location Based Services~~

Time Series and Historical Data Management

The above gives rise to *new research and development* in incorporating new data types, complex data structures, new operations and storage and indexing schemes in database systems.

# Extending Database Capabilities (2)

Background since the advent of the 21<sup>st</sup> Century:

First decade of the 21<sup>st</sup> century has seen tremendous growth in user generated data and automatically collected data from applications and search engines.

Social Media platforms such as Facebook and Twitter are generating millions of transactions a day and businesses are interested to tap into this data to “understand” the users

Cloud Storage and Backup is making unlimited amount of storage available to users and applications



# Extending Database Capabilities (3)

## Emergence of Big Data Technologies and NOSQL databases

New data storage, management and analysis technology was necessary to deal with the onslaught of data in petabytes a day ( $10^{15}$  bytes or 1000 terabytes) in some applications – this started being commonly called as “Big Data”.

Hadoop (which originated from Yahoo) and Mapreduce Programming approach to distributed data processing (which originated from Google) as well as the Google file system have given rise to Big Data technologies (Chapter 25). Further enhancements are taking place in the form of Spark based technology.

NOSQL (Not Only SQL- where SQL is the de facto standard language for relational DBMSs) systems have been designed for rapid search and retrieval from documents, processing of huge graphs occurring on social networks, and other forms of unstructured data with flexible models of transaction processing (Chapter 24).

# When not to use a DBMS

## Main inhibitors (costs) of using a DBMS:

- High initial investment and possible need for additional hardware.

- Overhead for providing generality, security, concurrency control, recovery, and integrity functions.

## When a DBMS may be unnecessary:

- If the database and applications are simple, well defined, and not expected to change.

- If access to data by multiple users is not required.

## When a DBMS may be infeasible:

- In embedded systems where a general purpose DBMS may not fit in available storage

Any questions?

# Late submission

8 late days you can use

20% reduction beyond this 8 days

Maximum of 2 late days in one submission

# Data Models

## Data Model:

A set of concepts to describe the ***structure*** of a database, the ***operations*** for manipulating these structures, and certain ***constraints*** that the database should obey.

## Data Model Structure and Constraints:

Constructs are used to define the database structure

Constructs typically include ***elements*** (and their ***data types***) as well as groups of elements (e.g. ***entity, record, table***), and ***relationships*** among such groups

Constraints specify some restrictions on valid data; these constraints must be enforced at all times

# Data Models (continued)

## Data Model Operations:

These operations are used for specifying database *retrievals* and *updates* by referring to the constructs of the data model.

Operations on the data model may include ***basic model operations*** (e.g. generic insert, delete, update) and ***user-defined operations*** (e.g. compute\_student\_gpa, update\_inventory)

# Categories of Data Models

## **Conceptual (high-level, semantic) data models:**

Provide concepts that are close to the way many users perceive data.

(Also called *entity-based* or *object-based* data models.)

## **Physical (low-level, internal) data models:**

Provide concepts that describe details of how data is stored in the computer. These are usually specified in an ad-hoc manner through DBMS design and administration manuals

## **Implementation (representational) data models:**

Provide concepts that fall between the above two, used by many commercial DBMS implementations (e.g. relational data models used in many commercial systems).

## **Self-Describing Data Models:**

Combine the description of data with the data values. Examples include XML, key-value stores and some NOSQL systems.

# Schemas versus Instances

## Database Schema:

The ***description*** of a database.

Includes descriptions of the database structure, data types, and the constraints on the database.

## Schema Diagram:

An ***illustrative*** display of (most aspects of) a database schema.

## Schema Construct:

A ***component*** of the schema or an object within the schema, e.g., STUDENT, COURSE.



# Schemas versus Instances

## Database State:

The actual data stored in a database at a ***particular moment in time***. This includes the collection of all the data in the database.

Also called database instance (or occurrence or snapshot).

The term *instance* is also applied to individual database components, e.g. *record instance*, *table instance*, *entity instance*

# Database Schema vs. Database State

## Database State:

Refers to the ***content*** of a database at a moment in time.

## Initial Database State:

Refers to the database state when it is initially loaded into the system.

## Valid State:


A state that satisfies the structure and constraints of the database.

## Distinction

The ***database schema*** changes very infrequently.

The ***database state*** changes every time the database is updated.

# Example of a Database Schema



## STUDENT

Name	Student_number	Class	Major
------	----------------	-------	-------

## COURSE

Course_name	Course_number	Credit_hours	Department
-------------	---------------	--------------	------------

## PREREQUISITE

Course_number	Prerequisite_number
---------------	---------------------

## SECTION

Section_identifier	Course_number	Semester	Year	Instructor
--------------------	---------------	----------	------	------------

## GRADE\_REPORT

Student_number	Section_identifier	Grade
----------------	--------------------	-------

**Figure 2.1**

Schema diagram for the database in Figure 1.2.

**COURSE**

Course_name	Course_number	Credit_hours	Department
Intro to Computer Science	CS1310	4	CS
Data Structures	CS3320	4	CS
Discrete Mathematics	MATH2410	3	MATH
Database	CS3380	3	CS

**SECTION**

Section_identifier	Course_number	Semester	Year	Instructor
85	MATH2410	Fall	04	King
92	CS1310	Fall	04	Anderson
102	CS3320	Spring	05	Knuth
112	MATH2410	Fall	05	Chang
119	CS1310	Fall	05	Anderson
135	CS3380	Fall	05	Stone

**GRADE\_REPORT**

Student_number	Section_identifier	Grade
17	112	B
17	119	C
8	85	A
8	92	A
8	102	B
8	135	A

**PREREQUISITE**

Course_number	Prerequisite_number
CS3380	CS3320
CS3380	MATH2410
CS3320	CS1310

**Figure 1.2**

A database that stores student and course information.

Example of a  
database  
state

# Three-Schema Architecture

Defines DBMS schemas at **three** levels:

**Internal schema** at the internal level to describe physical storage structures and access paths (e.g indexes).

Typically uses a **physical** data model.

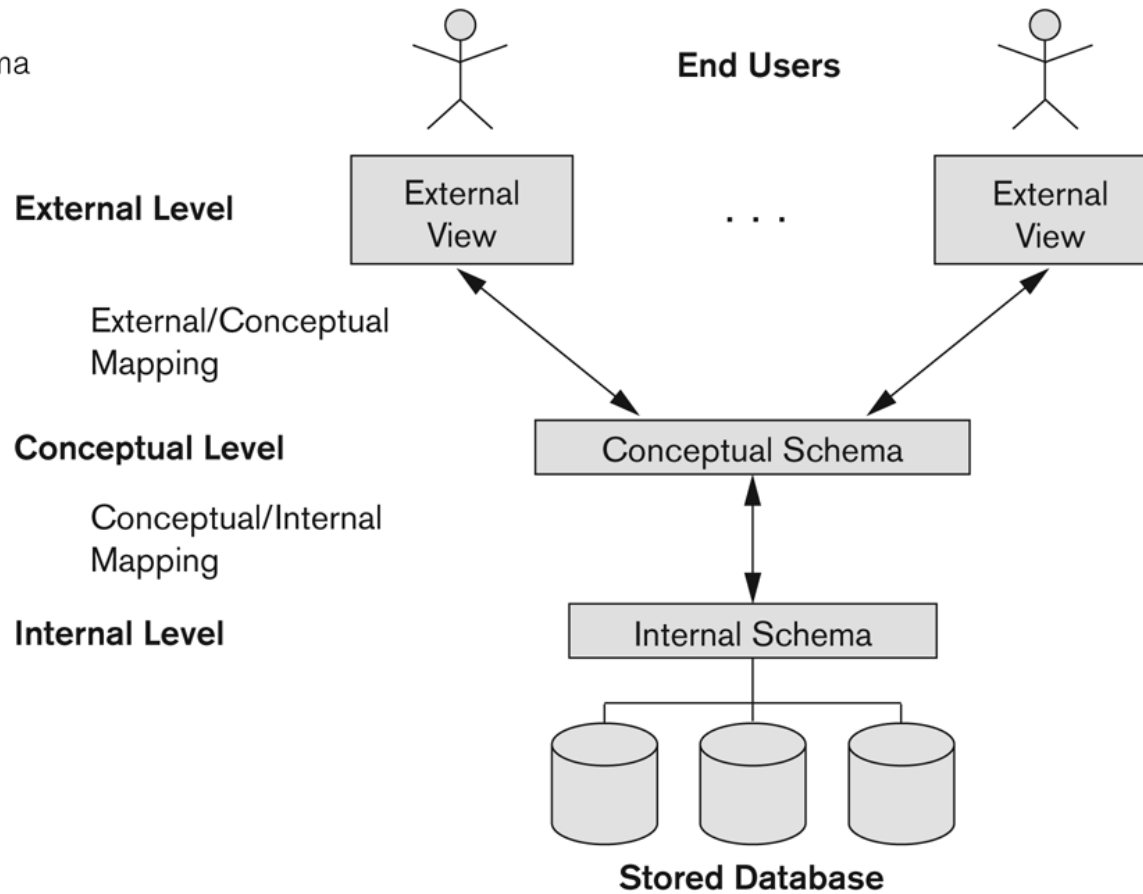
**Conceptual schema** at the conceptual level to describe the structure and constraints for the whole database for a community of users.

Uses a **conceptual** or an **implementation** data model.

**External schemas** at the external level to describe the various user views.

Usually uses the same data model as the conceptual schema.

**Figure 2.2**  
The three-schema  
architecture.



The three-  
schema  
architecture

# DBMS Languages

Data Definition Language (DDL)

Data Manipulation Language (DML)

# DBMS Languages

## **Data Definition Language (DDL):**

Used by the DBA and database designers to specify the conceptual schema of a database.

In many DBMSs, the DDL is also used to define internal and external schemas (views).

In some DBMSs, separate **storage definition language (SDL)** and **view definition language (VDL)** are used to define internal and external schemas.

SDL is typically realized via DBMS commands provided to the DBA and database designers



# DBMS Languages


## Data Manipulation Language (DML):

Used to specify database retrievals and updates

DML commands (data sublanguage) can be *embedded* in a general-purpose programming language (host language), such as COBOL, C, C++, or Java.

A library of functions can also be provided to access the DBMS from a programming language

Alternatively, stand-alone DML commands can be applied directly (called a *query language*).



# Types of DML

Select \* from  
Students.

## High Level or Non-procedural Language:

For example, the SQL relational language

Are “set”-oriented and specify what data to retrieve rather than how to retrieve it.

Also called **declarative** languages.

QBE – Query By Example

## Low Level or Procedural Language:

Retrieve data one record-at-a-time;

Constructs such as looping are needed to retrieve multiple records, along with positioning pointers.

COBOL / FORTRAN

# Bibliography / Acknowledgements

Instructor materials from Elmasri & Navathe 7e

 pk.profgiri

 Ponnurangam.kumaraguru

 /in/ponguru

 ponguru

 pk.guru@iiit.ac.in

Thank you  
for attending  
the class!!!